

# Thoth: the Python cloud resolver

Thoth team - [thoth-station.ninja](https://thoth-station.ninja)

# Agenda

1. \$ whoeare
2. Introducing the Python cloud resolver
3. Using Python cloud resolver
4. Declarative interface for the resolver - prescriptions
5. Containerized Python applications
6. Security - AIDevSecOps
7. Cross-index resolution
8. References

## \$ whoarewe

- Thoth - AIDevSecOps
  - Started (2018) as a research project in Red Hat AI CoE team, Office of the CTO
  - [thoth-station.ninja](https://thoth-station.ninja)
- See our linked [YouTube channel](#) for more information
- Follow us on Twitter - [@ThothStation](#)

# Our mission

- Help Python developers and data scientists create healthy applications
- Project has multiple parts:
  - [AICoE-CI](#) - a CI that builds container images
  - [Thoth resolver](#) - a recommendation engine for Python applications
    - [AIDevSecOps](#)
  - [Dependency Monkey](#) - a service that can validate software in a cluster
  - [jupyterlab-requirements](#) extension for managing dependencies
  - [Bots maintaining GitHub repositories](#)
  - [A self hosted Python package index using Pulp](#) available to all Red Hatters
  - [Container image analysis and containerized Python applications](#)



## *Introducing the Python cloud resolver*



## Python resolvers

- pip
  - the package installer for Python
- Pipenv
  - Python development workflow for humans
- Poetry
  - Python dependency management and packaging made easy
- Thoth
  - Resurrected ancient deities helping humans with software development



## Python resolvers

- pip
  - the package installer for Python
- Pipenv
  - Python development workflow for humans
- Poetry
  - Python dependency management and packaging made easy
- Thoth
  - Resurrected ancient deities helping humans with software development

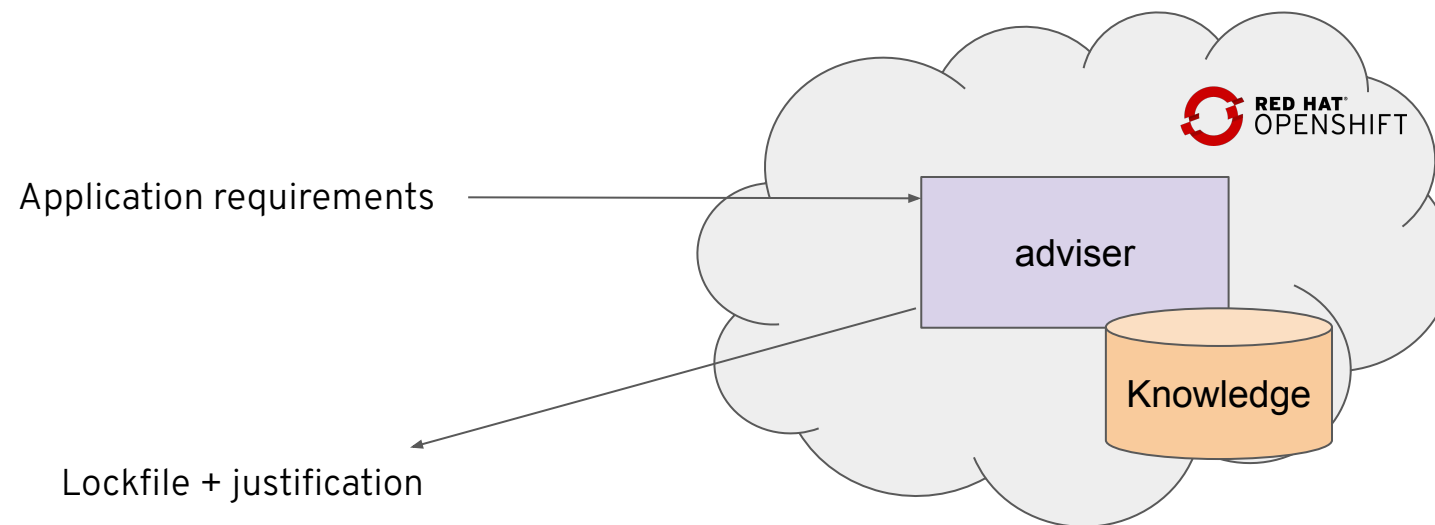
*Latest software is not always the greatest choice.*

## The Python resolver run in cloud

- Recommendation engine for Python applications
- Publicly available to the community
- Stochastic resolver implementing gradient-free reinforcement learning methods
  - Temporal difference learning is used in production
- See documentation for more information:
  - [thoth-station.ninja/docs/developers/adviser](https://thoth-station.ninja/docs/developers/adviser)



# Python cloud resolver



# Python cloud resolver

Requirements

Constraints

Information about software in runtime environment

- OS, Python version, base image, CUDA, ...

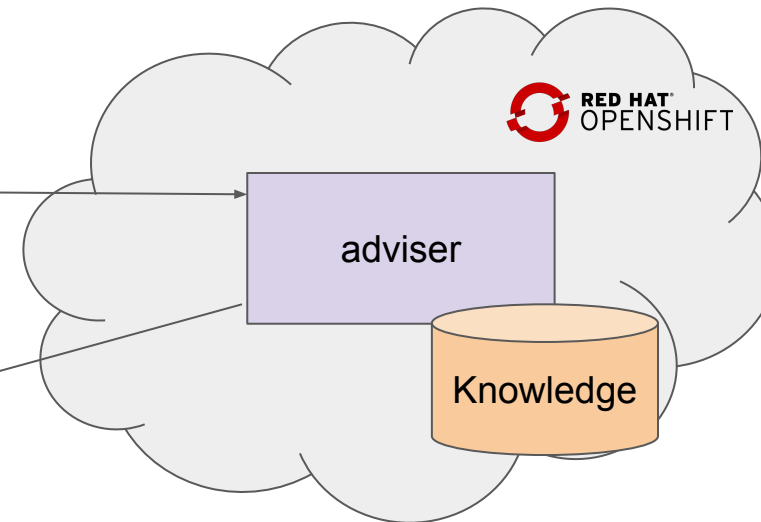
Information about hardware in runtime environment

- CPU, GPU, ...

Static source-code analysis

Recommendation type

Lockfile + justification



## *Using Python cloud resolver*



[Managing security in Python applications with the Thoth cloud Python resolver](#)

# Thamos Command Line Interface

- One of the Thoth client tools, other tools:
  - [jupyterlab-requirements](#)
  - [Kebechet bot](#)
- Talks to Thoth's backend and helps with managing your environment
- Available on PyPI:

```
$ pip install thamos
$ thamos --help
$ thamos config
```

# Runtime environments and recommendation types

- Concept of “*overlays*”
  - Each overlay states requirements for the application for the given runtime environment
  - Ex. runtime environment for inference and runtime environment for training
    - Can have different requirements and runtime environments
- Support for “*recommendation types*”
  - Different resolution considering intention with the application
  - [thoth-station.ninja/recommendation-types](https://thoth-station.ninja/recommendation-types)
  - Ex. production environment should be secure, isolated environment where the training is done should be performant

```

host: khemenu.thoth-station.ninja
requirements_format: pipenv

runtime_environments:
- name: "inference"
  recommendation_type: security
  operating_system:
    name: "rhel"
    version: "8"
  python_version: "3.8"
- name: "training"
  recommendation_type: performance
  operating_system:
    name: "rhel"
    version: "8"
  cuda_version: "11.1"
  python_version: "3.8"

```

```

[[source]]
url = "https://pypi.org/simple"
verify_ssl = true
name = "pypi-org-simple"

[packages]
flask = "*"
tensorflow = "~=2.4"

[dev-packages]

[requires]
python_version = "3.8"

[thoth]
disable_index_adjustment = false

[thoth.allow_prereleases]

```

```

[[source]]
url = "https://pypi.org/simple"
verify_ssl = true
name = "pypi-org-simple"

[[source]]
url = "https://thoth-station.ninja/simple"
verify_ssl = true
name = "aicoe"

[packages]
boto3 = "*"
tensorflow = {"version"="~=2.4", "index"="aicoe"}

[dev-packages]

[requires]
python_version = "3.8"

[thoth]
disable_index_adjustment = true

[thoth.allow_prereleases]
protobuf = true

```

## *Demo: Thamos CLI*

*Declarative interface for the resolver to resolve Python packages following prescribed rules*

 [Thoth prescriptions for resolving Python dependencies](#)



# Resolution pipeline

Requirements

Constraints

Information about software in runtime environment

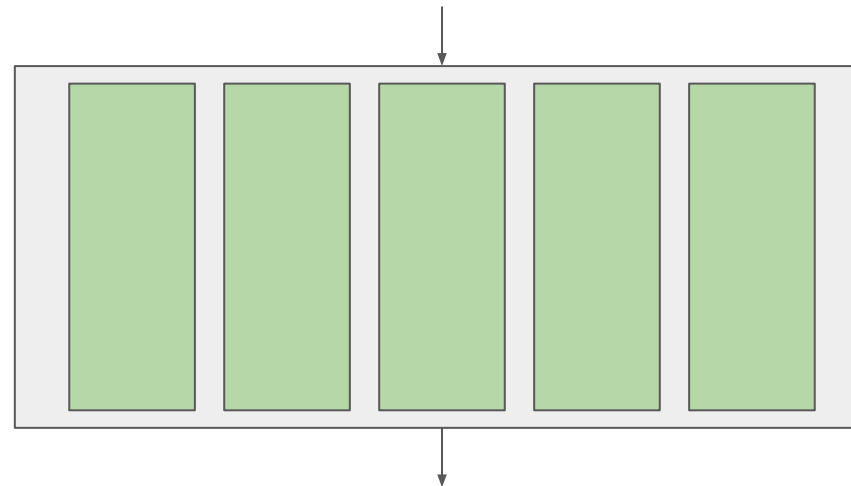
- OS, Python version, base image, CUDA, ...

Information about hardware in runtime environment

- CPU, GPU, ...

Static source-code analysis

Recommendation type



Lockfile + justification

# Resolution pipeline

- The whole resolution process is seen as a pipeline
- The resolution pipeline is made out of pipeline units
- Each pipeline unit has its own type
- Pipeline units can be implemented in Python or declared in YAML files

## Prescriptions - declarative interface to the cloud based resolver

- Provide a way to declaratively state how the resolution process should look like
- Community driven open database used by the resolver to resolve high quality software
  - [github.com/thoth-station/prescriptions](https://github.com/thoth-station/prescriptions)
- A set of YAML files that are automatically consumed by the resolver in a deployment
- See documentation for more information:
  - [thoth-station.ninja/docs/developers/adviser/prescription.html](https://thoth-station.ninja/docs/developers/adviser/prescription.html)

# Prescriptions

*When using OpenShift or Kubernetes, one provides manifest files that state how the desired state of a cluster should look like. Prescriptions might be seen analogous to this - prescriptions provide a way to declaratively state how the desired dependency resolution should look like considering the prescribed rules. Then, it's up to the reinforcement learning algorithm implemented in Thoth's adviser to find a solution in the form of a lockfile respecting the prescribed rules, requirements for the application and other inputs to the Thoth's cloud resolver.*

# Prescriptions - Example

- Pillow in version 8.3.0 does not work with NumPy

[github.com/python-pillow/Pillow/issues/5571](https://github.com/python-pillow/Pillow/issues/5571)

```
with PIL.Image.open(filepath) as img:  
    numpy.array(img, dtype=numpy.float32)
```

```
> frame_paletted = np.array(im, np.uint8)  
E TypeError: __array__() takes 1 positional argument but 2 were given
```

```
/lib/python3.9/site-packages/imageio/plugins/pillow.py:745: TypeError
```

```
units:
steps:
- name: Pillow830TypeErrorStep
  type: step
  should_include:
    adviser_pipeline: true
  match:
    package_version:
      name: pillow
      version: ==8.3.0
      index_url: https://pypi.org/simple
    state:
      resolved_dependencies:
        - name: numpy
run:
  not_acceptable: Pillow in version 8.3.0 does not work with NumPy
  stack_info:
    - type: WARNING
      message: Pillow in version 8.3.0 does not work with NumPy
      link: https://github.com/python-pillow/Pillow/issues/5571
```

units:

steps:

- name: Pillow830TypeErrorStep

type: step

should\_include:

- adviser\_pipeline: true

match:

package\_version:

- name: pillow

- version: ==8.3.0

- index\_url: <https://pypi.org/simple>

state:

resolved\_dependencies:

- name: numpy

run:

not\_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack\_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:

steps:

```
- name: Pillow830TypeErrorStep  
  type: step
```

should\_include:

adviser\_pipeline: true

match:

package\_version:

name: pillow

version: ==8.3.0

index\_url: https://pypi.org/simple

state:

resolved\_dependencies:

- name: numpy

run:

not\_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack\_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>



units:

steps:

- name: Pillow830TypeErrorStep

type: step

should\_include:

adviser\_pipeline: true

match:

package\_version:

name: pillow

version: ==8.3.0

index\_url: https://pypi.org/simple

state:

resolved\_dependencies:

- name: numpy

run:

not\_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack\_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:

steps:

- name: Pillow830TypeErrorStep

type: step

should\_include:

adviser\_pipeline: true

match:

package\_version:

name: pillow

version: ==8.3.0

index\_url: https://pypi.org/simple

state:

resolved\_dependencies:

- name: numpy

run:

not\_acceptable: Pillow in version 8.3.0 does not work with NumPy

stack\_info:

- type: WARNING

message: Pillow in version 8.3.0 does not work with NumPy

link: <https://github.com/python-pillow/Pillow/issues/5571>

units:  
steps:  
- name: Pillow830TypeErrorStep  
type: step  
should\_include:  
  adviser\_pipeline: true  
match:  
  package\_version:  
    name: pillow  
    version: ==8.3.0  
    index\_url: https://pypi.org/simple  
state:  
  resolved\_dependencies:  
    - name: numpy

run:  
  not\_acceptable: Pillow in version 8.3.0 does not work with NumPy  
  stack\_info:  
    - type: WARNING  
      message: Pillow in version 8.3.0 does not work with NumPy  
      link: <https://github.com/python-pillow/Pillow/issues/5571>



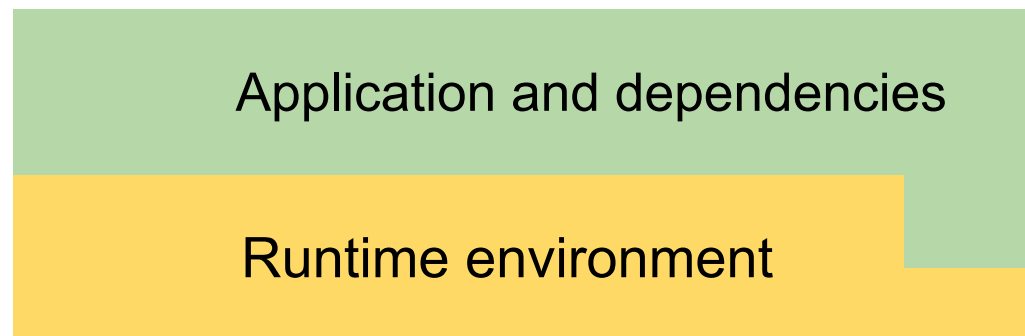
## *Demo: Prescriptions*

## *Containerized Python applications*

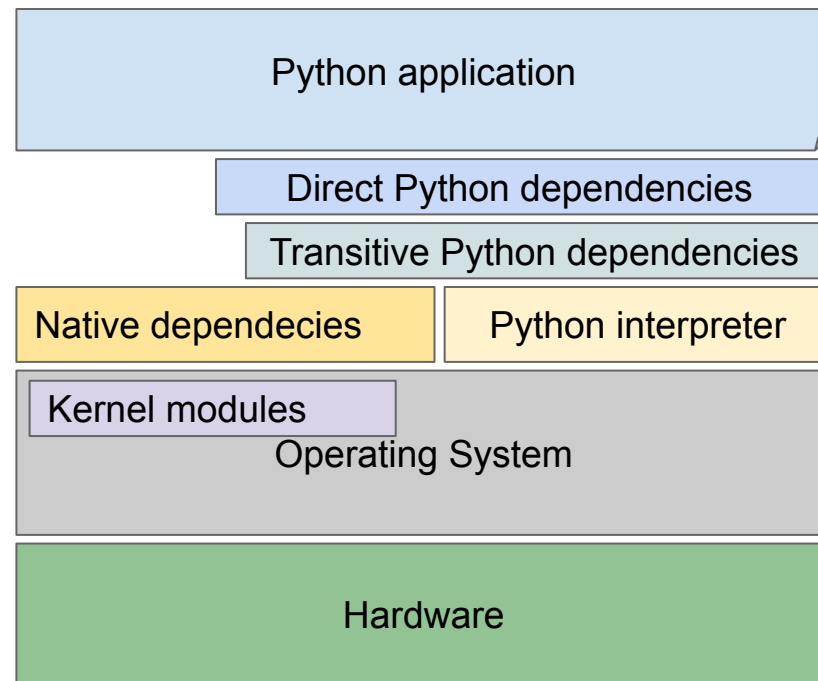
 [Build and extend containerized applications with Project Thoth](#)

# Fitting the containerized runtime environment

- RPM, Python packages, ABI, CUDA, cuDNN, OpenMKL, ...
- Resolver considers containerized environment
  - Still optional, the resolution might be “generic”, like pip, Pipenv, ...



# Fitting the runtime environment

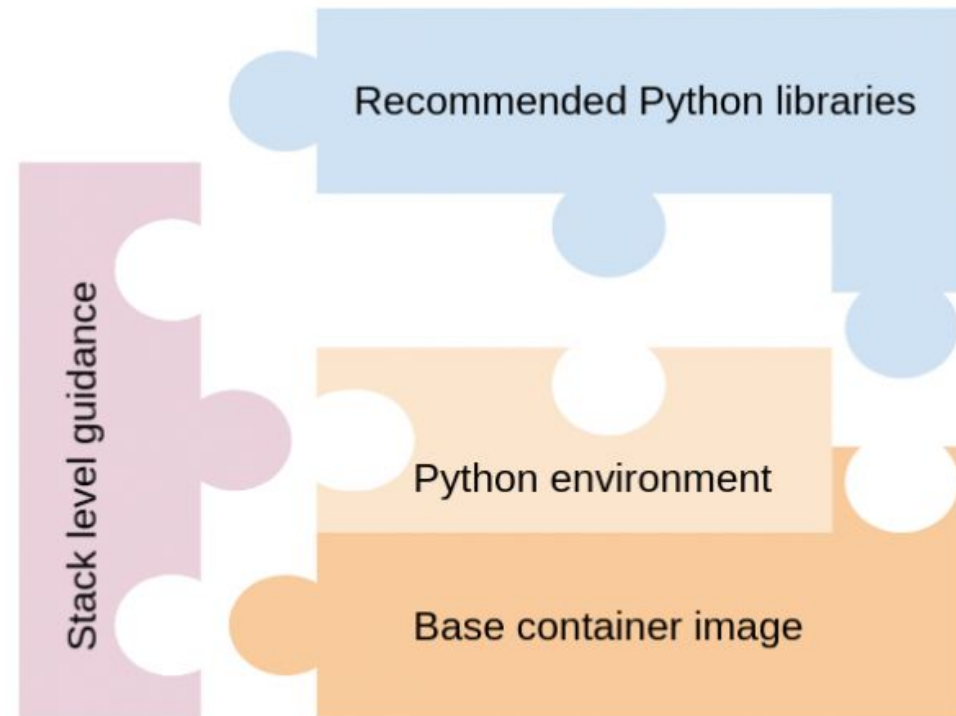


# Fitting the containerized runtime environment

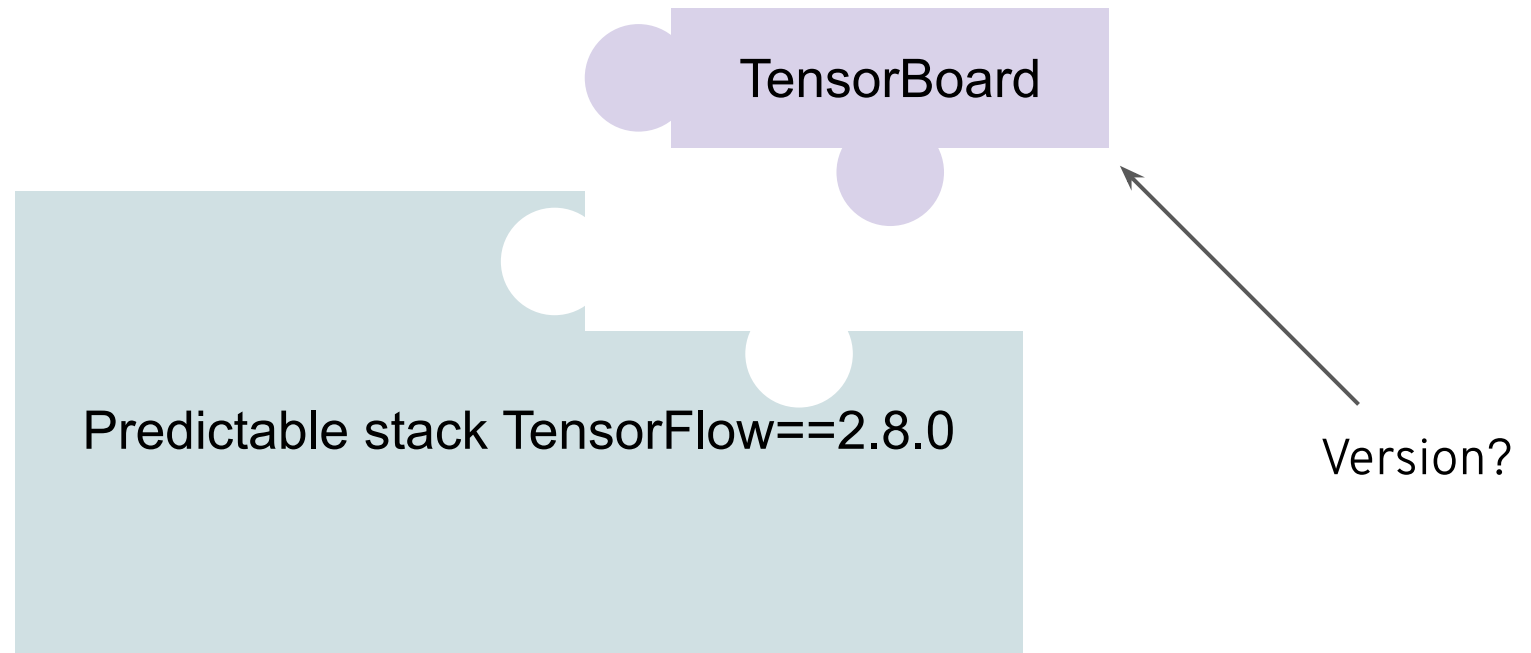
- Base container image
  - Python base container images
  - Predictable stacks
  - JupyterLab container images
  - ...
- Synthesize Python layer and base container image layer



## Fitting the environment example



## Fitting the environment example



# thamos images --help

Usage: thamos images [OPTIONS]

Check available Thoth container images.

Examples:



```
thamos images --output-format json
```

```
thamos images --os-name fedora --os-version 35 --python-version 3.9
```

```
thamos images --symbol GLIBC_FOO
```

Options:

- o, --output-format [json|yaml|table]  
Specify output format for the status report.
- n, --os-name OS\_NAME Operating system name filter.
- v, --os-version OS\_VERSION Operating system version filter.
- p, --python-version PY\_VERSION  
Python interpreter version filter.
- cuda-version CUDA\_VERSION CUDA version filter.
- image-name IMAGE\_NAME Filter based on image name.
- library-name LIBRARY\_NAME Filter based on library name.
- symbol SYMBOL Filter based on symbol.
- package-name PACKAGE\_NAME Filter based on Python package name.
- rpm-package-name RPM\_PACKAGE\_NAME  
Filter based on RPM package name.
- help Show this message and exit.



## *Demo: Thamos and container images*

## *Security - AIDevSecOps*



[Secure your Python applications with Thoth recommendations](#)

# Security - AIDevSecOps

- Docs: [Thoth security advises](#)
  - Recommendations based on static source code analysis
    - [See recommendations from the Python standard library \(example\)](#)
  - PyPA - advisory-db
    - A database of known vulnerabilities in Python ecosystem
    - [github.com/pypa/advisory-db](https://github.com/pypa/advisory-db)
  - Security Scorecards by Open Source Security Foundation
    - [openssf.org/blog/2020/11/06/security-scorecards-for-open-source-projects](https://openssf.org/blog/2020/11/06/security-scorecards-for-open-source-projects)
    - Example: see [scorecards prefixed prescriptions for TensorFlow](#)
  - Container image analyses - vulnerabilities in the base container images used
- ... additional information about Python packages not strictly related to security*



## *Demo: Securing Python applications*

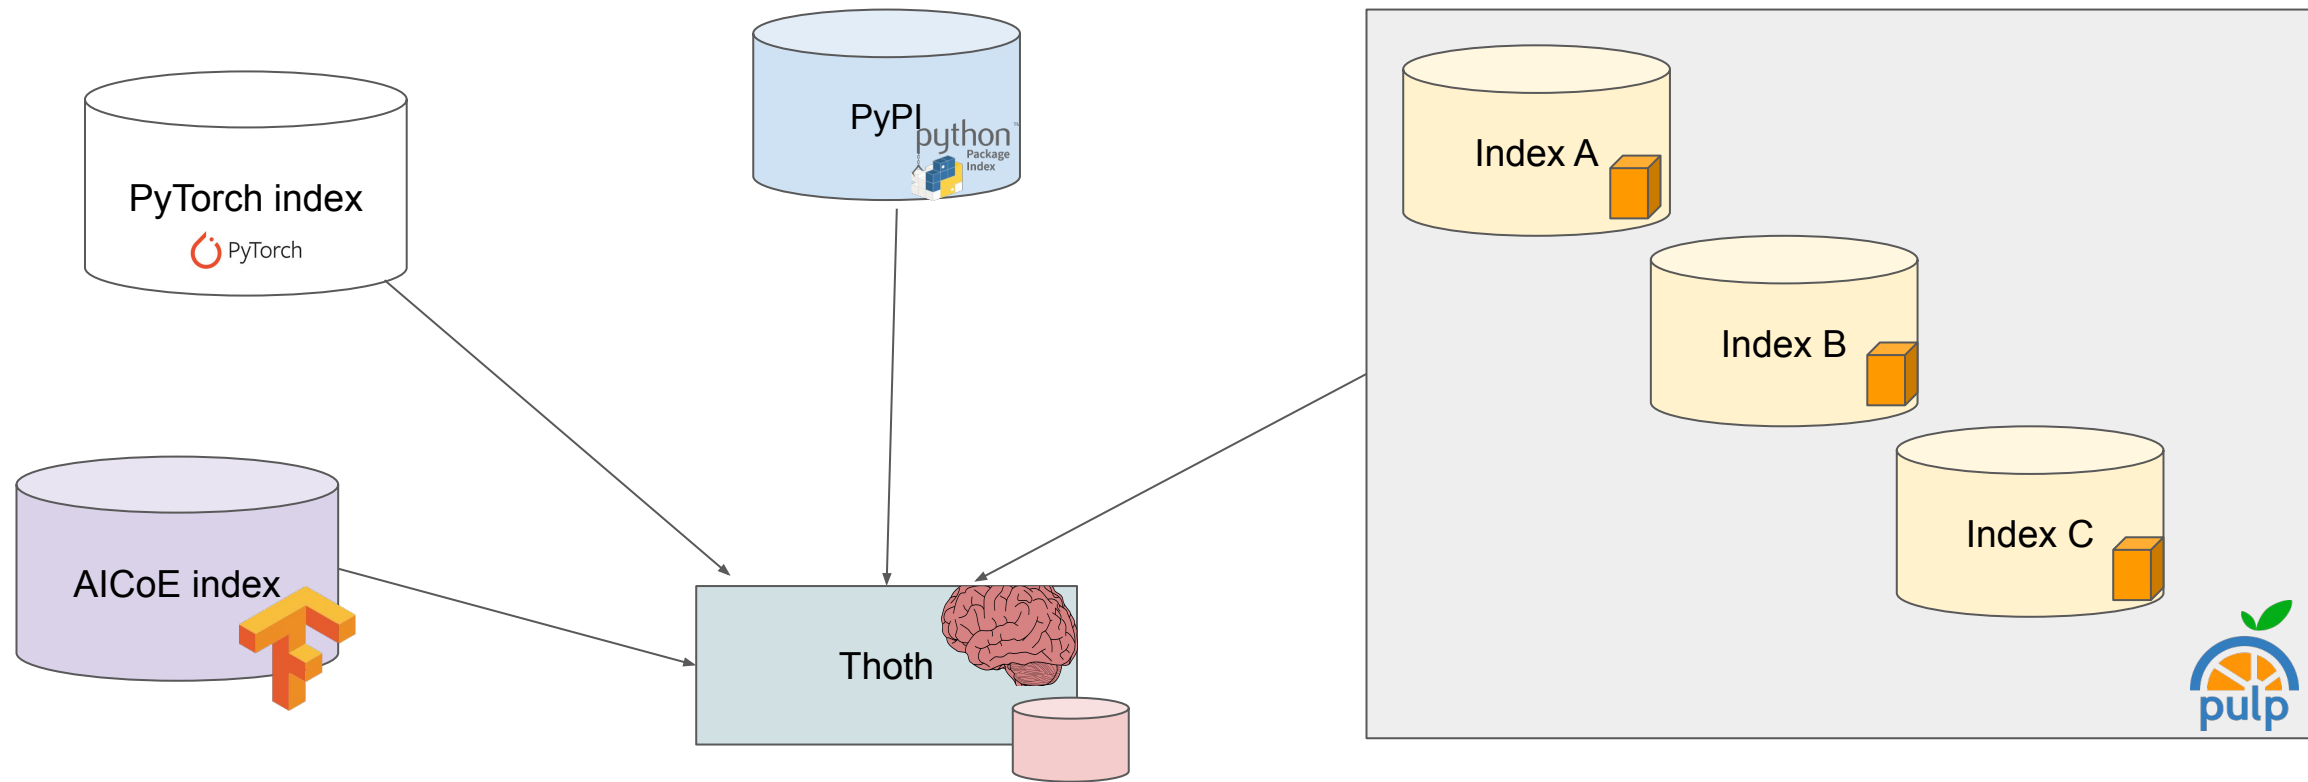
*Cross-index resolution*



[Prevent Python dependency confusion attacks with Thoth](#)



# Python package indexes & Thoth



# Python package indexes

- Automatically monitored in a Thoth deployment
  - e.g. builds outside of manylinux standards:
    - index with CUDA specific builds
    - AVX2 optimized builds of TensorFlow
- Packages published on monitored Python package indexes are automatically analyzed
- Cross index resolution without dependency confusion
- One central point of knowledge in cloud based resolver

 *Demo: Thamos and cross-index resolution*

## *References*



# References

[YouTube channel](#) [News](#) [Talks](#) [Datasets](#) [Documentation](#) [Package index](#) [API](#) [Status](#) [Tutorial](#)

[Get involved](#)

## Project Thoth

Using Artificial Intelligence to analyse and recommend software stacks for Python applications.

[Get started](#)



[thoth-station.ninja](https://thoth-station.ninja)



## References

- [Introspecting containerized Python applications in a cluster with Thoth Amun](#)
- [How to self-host a Python package index using Pulp](#)
- [Extracting dependencies from Python packages](#)
- [Extracting information from Python source code](#)
- [Prevent Python dependency confusion attacks with Thoth](#)
- [Build and extend containerized applications with Project Thoth](#)
- [Customize Python dependency resolution with machine learning](#)
- [Generating pseudorandom numbers in Python](#)
- [Secure your Python applications with Thoth recommendations](#)
- [Find and compare Python libraries with project2vec](#)

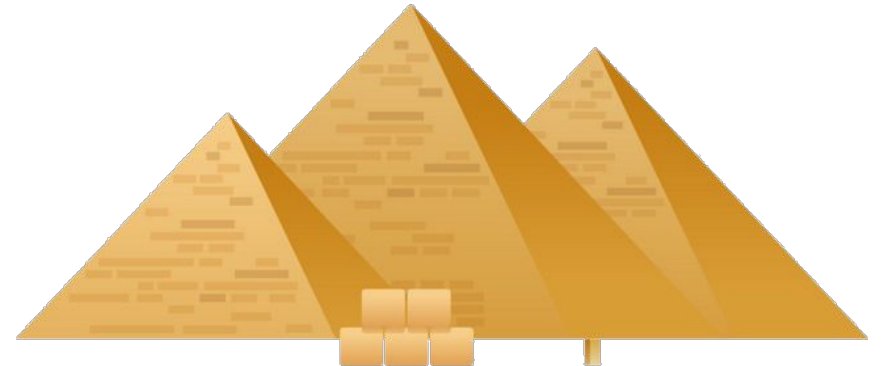


## References

- [Thoth prescriptions for resolving Python dependencies](#)
- [Resolve Python dependencies with Thoth Dependency Monkey](#)
- [micropipenv: Installing Python dependencies in containerized applications](#)
- [Continuous learning in Project Thoth using Kafka and Argo](#)
- [Can we consider --editable a bad practice?](#)
- [Managing Python dependencies with the Thoth JupyterLab extension](#)
- [Use Kebechet machine learning to perform source code operations](#)
- [AI software stack inspection with Thoth and TensorFlow](#)
- [Microbenchmarks for AI applications using Red Hat OpenShift on PSI in project Thoth](#)

## References

- Thoth's website
  - [thoth-station.ninja](https://thoth-station.ninja)
- Source code:
  - [github.com/thoth-station](https://github.com/thoth-station)
- [@ThothStation](#) Twitter handle
- [Thoth Station YouTube channel](#)
- [Talks and presentations](#)





# Thank you

Red Hat is the world's leading provider of enterprise open source software solutions. Award-winning support, training, and consulting services make Red Hat a trusted adviser to the Fortune 500.



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)